

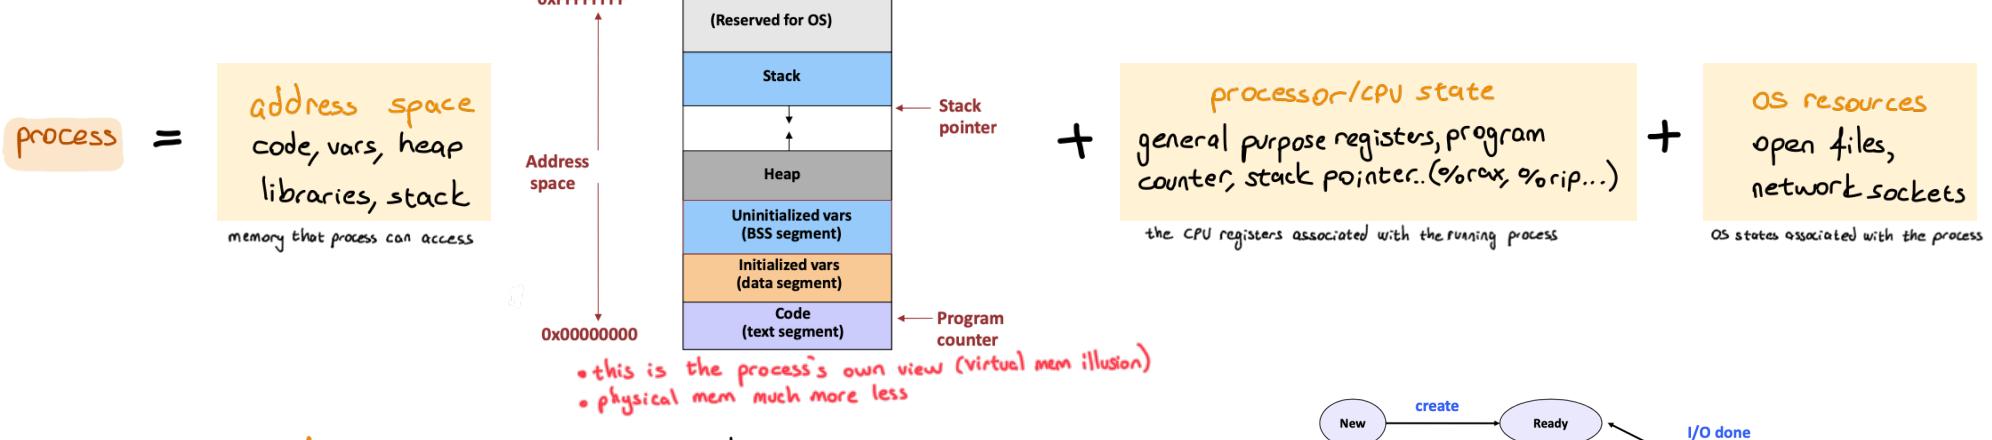
2-processes

program = countable set of ordered and sequential instructions that needs to be executed to accomplish a task

process = execution part of a program, provides programs two abstractions (CPU and mem)

① logical control flow = illusion of having own CPU → with interleaved executions of processes (multitasking)

② private virtual address = illusion of having own main memory → with managing address space by virtual memory system



execution states of a process = running = currently using CPU

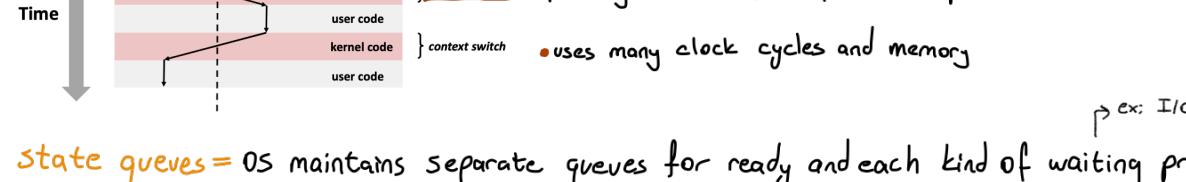
↳ ready = waiting to be assigned to CPU

↳ waiting/sleeping = waiting for an event (I/O, timer)

process control block = OS maintains a PCB for each process

↳ process id + user id + execution state + saved CPU state + OS resources + mem management info + scheduling priority + accounting information

kernel = central component of OS that manages operations of computer and hardware → processes are managed by kernel



state queues = OS maintains separate queues for ready and each kind of waiting processes

↳ PCBs move between these queues, when ready processes queue is completed, it moves to waiting queue to ready queue

concurrent processes = if two processes' flows overlap in time

↳ otherwise they are sequential

creating processes

↳ parent process creates a new running child process by calling fork(system call), they run concurrently

int fork(void) = returns 0 to child process

↳ returns child's pid to parent process (returns -1 if failed) pid = process id

• child gets an identical but separate copy of the parent's virtual address space and parent's open file descriptors → copy of address space + CPU state + OS resources

• called ones but returns twice

```
int main(int argc, char** argv)
{
    pid_t pid;
    int x = 1;

    pid = fork();
    if (pid == 0) { /* Child */
        printf("child : %d\n", ++x);
        main_fork();
        exit(0);
    }
    /* Parent */
    printf("parent: %d\n", --x);
    main_fork();
    return 0;
}
```

```
int main(int argc, char** argv)
{
    pid_t pid;
    int x = 1;

    pid = fork();
    if (pid == 0) { /* Child */
        printf("child : %d\n", ++x);
        main_fork();
        exit(0);
    }
    /* Parent */
    printf("parent: %d\n", --x);
    main_fork();
    return 0;
}
```

```
void fork()
{
    printf("1\n");
    fork();
    printf("2\n");
    fork();
    printf("3\n");
}

void fork()
{
    printf("1\n");
    fork();
    printf("2\n");
    fork();
    printf("3\n");
}
```

```
void fork()
{
    printf("1\n");
    fork();
    printf("2\n");
    fork();
    printf("3\n");
}

void fork()
{
    printf("1\n");
    fork();
    printf("2\n");
    fork();
    printf("3\n");
}
```

↳ In graphs parents are always at bottom

terminating processes → calling the exit func → void exit(int status) → status=0 normal, status!=0 error → exit(0); called once but never returns → because function that called is no longer exist

↳ returning from the main

↳ receiving a signal whose default action is to terminate ctrl+c

atexit(func-name) = the func pointed by atexit is automatically called without arguments when the program terminates normally

↳ returns zero if the function registration is successful, nonzero else

• when process terminates it still consumes system resources (exit status, OS tables) → zombie

reaping child processes performed by parent on terminated child using wait or waitpid, kernel then deletes child.

↳ if parent terminates without reaping a child, the child will be reaped by init process (pid=1)

↳ so only need explicit reaping in long-running processes → shells, servers

int wait(status *child_status) = suspends current parent process until one of its children terminates

• return value is pid of the child process that terminated

• if child_status != NULL, then the int it points to will be set to a val that indicates reason the child terminated and the exit status

process completion status

```
void fork()
{
    int child_status;
    if (fork() == 0) {
        printf("HC: hello from child\n");
        exit(0);
    }
    else {
        printf("HP: hello from parent\n");
        wait(&child_status);
        printf("CT: child has terminated\n");
    }
    printf("Bye\n");
}
```

```
void fork()
{
    pid_t child_status;
    if (fork() == 0) {
        for (int i = 0; i < N; i++)
            if (pid == 0) { /* Child */
                printf("1\n");
                fork();
                printf("2\n");
                fork();
                printf("3\n");
            }
        if (WIFEXITED(child_status)) {
            printf("Child %d terminated with exit status %d\n",
                pid, WEXITSTATUS(child_status));
        }
        else {
            printf("Child %d terminated abnormally\n", pid);
        }
    }
}
```

↳ all children exits arbitrary, but they are zombie until parent is terminated

↳ parent access them in order

loading and running programs

loads and runs in the current process

int execve(char *filename, char *argv[], char *environ[]) → base function called ones but never returns, if no error

• argv[0] = filename by convention, in the form "name=value" ex: user=dash putenv, getenv, printenv

• Overwrites code, data, and stack → retains pid, open files and signal context

• does not fork a new process, it replaces the address space and CPU state of the current process

• loads the new address space from the executable file and starts it from main()

* so start a new program, use fork() followed by execve()

int execve(char *path, char *arg0, char *arg1, ..., 0)

• environment taken from char **environ

environment variable = dynamic named val that can affect the way running process will behave on a computer

exec() func family

The suffix's determine the arguments

- 1: arguments are passed as a list of strings to the main()
- v: arguments are passed as an array of strings to the main()
- p: used to search for the new running program
- e: the environment can be specified by the caller

One can mix them with different combinations

- int execve(const char *path, const char *arg, ...);
- int execve(const char *file, const char *arg, ...);
- int execve(const char *path, const char *const arg[]);
- int execve(const char *path, char *const argv[]);
- int execve(const char *path, char *const argv[]);
- int execve(const char *file, char *const argv[]);
- int execve(const char *file, char *const argv[]);

The initial argument is always the name of a file to be executed.

```
int main(int argc, char **argv) {
    if (argc > 1) { /* Parent process */
        if (fork() == 0) { /* Child process */
            char args[3];
            if (args[0] == 'a') /* Child has args */
                args[1] = 'a';
            else {
                args[1] = 'b';
                args[2] = '\0';
            }
            if (execve("./bin/echo", args) == -1) {
                perror("Warning: execve returned an error.\n");
                exit(-1);
            }
            printf("Child process should never get here\n");
            exit(42);
        }
    }
}
```