

## 4- file abstraction

virtual filesystem (VFS) interface = file is the abstraction used for I/O

↳ used to access disks, CDs, DVDs, USBs, serial devices, network → /dev/usb/l, /dev/mouse, /dev/kbd <sup>keyboard</sup>

C standard I/O = buffering between program and actual files

↳ fopen, fread, fscanf, fprintf → not system calls by themselves

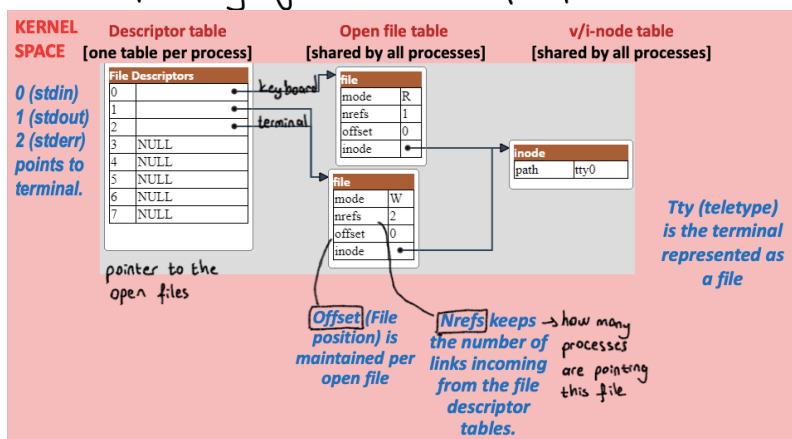
UNIX I/O = open, read, write, stat, close, lseek → changing the current file position (offset in file)

- file is a sequence of bytes → all I/O devices represented as files (even kernel)

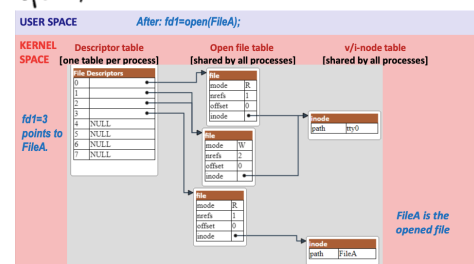
opening files = `int file-descriptor = open(" /etc/hosts", 0_RDONLY)` if == -1 → error

\* every process has three special files already open = `stdin`, `stdout`, `stderr` → standard error

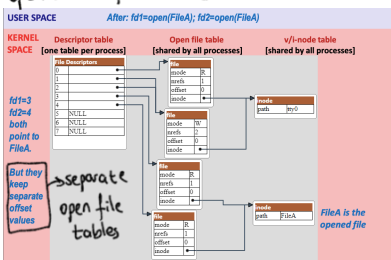
operating system resources per process



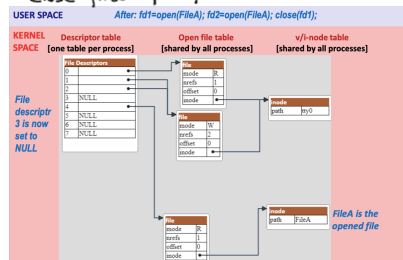
open file1



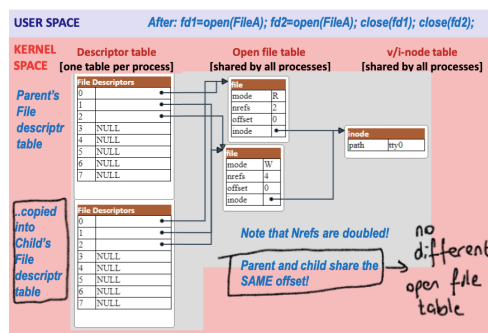

open same file1 twice



close first open file1



close second open file2



after fork():

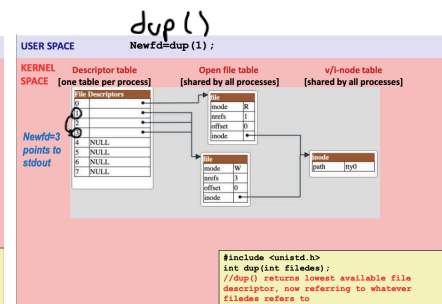
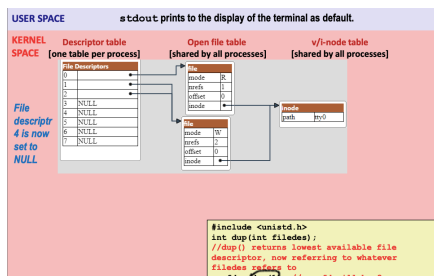
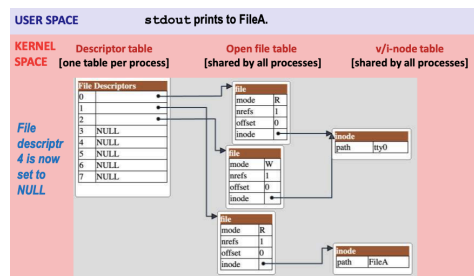
way to communicate between child and parent

shell redirection

> ./myprogram > somefile.txt → writes output to somefile (connects stdout to somefile.txt)

> ./myprogram < input.txt > somefile.txt → connects stdin to input.txt, stdout to somefile.txt

> ./myprogram 2 > error.txt → connects stderr to error.txt



↳ connect stdout to file A, can't be done directly because it is kernel space

- system calls are needed

dup() system call

↳ duplicates fd1

↳ returns new index

dup(fd)

```
//Descriptor table
void *DT[maxFd];

int dup(int oldfd){
    //get the lowest available
    //File descriptor
    newfd = lowestFD(DT);
    DT[newfd]=DT[oldfd];
    return(newfd);
}
```

copies the pointer of file table

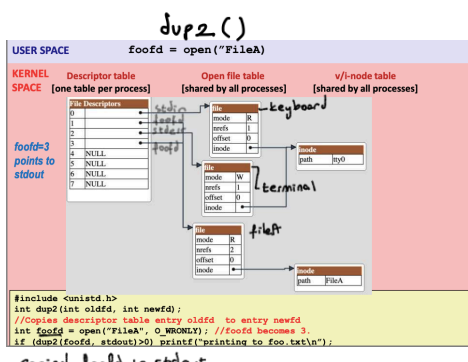
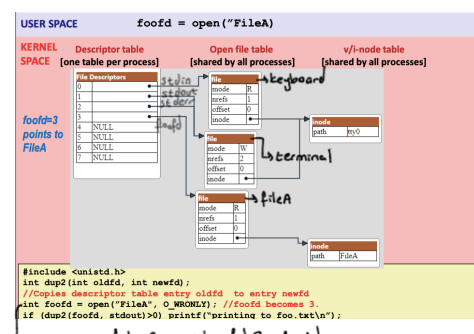
dup2(oldfd, newfd)

```
//Descriptor table
void *DT[maxFd];

int dup2(int oldfd, int newfd){
    DT[newfd]=DT[oldfd];
    return(newfd);
}
```


copies descriptor table entry oldfd to entry newfd

- if oldfd not valid → fail
- ↳ newfd still open




example-1

```
#include "csapp.h"
int main(int argc, char *argv[])
{
    int fd1, fd2, fd3;
    char c1, c2, c3;
    char *fname = argv[1];
    fd1 = open(fname, O_RDONLY, 0); //4
    fd2 = open(fname, O_RDONLY, 0); //5
    fd3 = open(fname, O_RDONLY, 0); //6
    dup2(fd2, fd3); //they share same file offset
    read(fd1, &c1, 1); //a
    read(fd2, &c2, 1); //a
    read(fd3, &c3, 1); //a
    printf("c1 = %c, c2 = %c, c3 = %c\n", c1, c2, c3);
    return 0; //file contains "abcde"
}
```



example-2

```
#include "csapp.h"
int main(int argc, char *argv[])
{
    int fd1;
    int s = getpid() & 0x1;
    char c1, c2;
    char *fname = argv[1];
    fd1 = open(fname, O_RDONLY, 0);
    Read(fd1, &c1, 1); //a
    if (fork()) { /* Parent */
        sleep(s);
        Read(fd1, &c2, 1); //a
        printf("Parent: c1 = %c, c2 = %c\n", c1, c2);
    } else { /* Child */
        sleep(1-s);
        Read(fd1, &c2, 1); //a
        printf("Child: c1 = %c, c2 = %c\n", c1, c2);
    }
    return 0; //file contains "abcde"
}
```




buffering in standar I/O

buf → hello\n ... → output if "\n" or fflush() is seen, no separate system calls

strace ./hello → prints a list of system calls made by the program, I write syscall

```
void fork2()
{
    printf("LO\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```



```
void fork2a()
{
    printf("LO\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```

2-20

- LO printed twice, because it is buffered without "\n" and buffer is duplicated to child