

threads

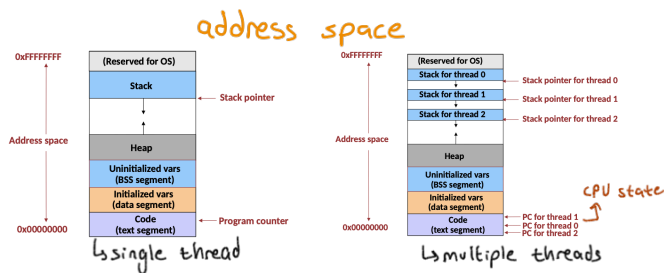
thread = unit of CPU scheduling

↳ each process has one or more threads within it

↳ each thread has its own stack, CPU registers, program counter (CPU execution state)

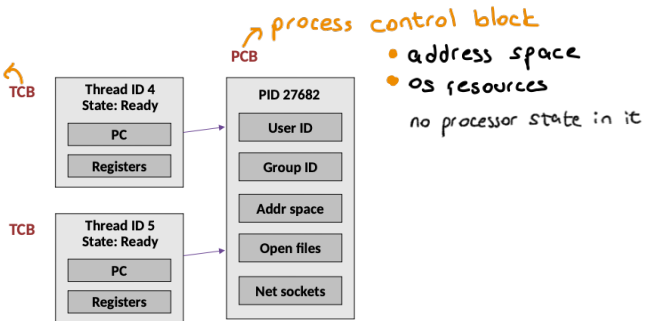
↳ all threads within a same process share the same address space, OS resources

process = address space + OS resources + CPU state




thread control block

- processor state
- pointer to PCB



context switching = copy TCB to PCB from ready queue, no need to change address space (switching CPU state)

user level thread address space = you allocate space for stacks in heap



int setjmp(jmp_buf env) = save CPU state

↳ 0 if saving the current state

↳ returns rtnval if call is from longjmp

void longjmp(jmp_buf env, int rtnval) = restore

struct jmp_buf {} → CPU specific fields

- can only jump to func that has not been completed yet →

```
int main(int argc, void *argv) {
    int i, restored = 0;
    jmp_buf saved;

    for (i = 0; i < 10; i++) {
        printf("Value of i is now %d\n", i);
        if (i == 5) {
            printf("OK, saving state...\n");
            if (setjmp(saved) == 0) {
                printf("Saved CPU state and breaking from loop.\n");
                break;
            } else {
                printf("Restored CPU state, continuing where we saved\n");
                restored = 1;
            }
        }
    }
    if (!restored) longjmp(saved, 1);
}
```

① returns 0 to save


② call for restore

③ returns 1 to restore

④ exit

output

```
Value of i is now 0
Value of i is now 1
Value of i is now 2
Value of i is now 3
Value of i is now 4
Value of i is now 5
OK, saving state...
Saved CPU state and breaking from loop.
Restored CPU state, continuing where we saved
Value of i is now 6
Value of i is now 7
Value of i is now 8
Value of i is now 9
```



preemption = forcefully taking CPU from a thread, and give it to someone else

↳ nonpreemptive threads = CPU allocates the process till it terminates, can be infinitely long

↳ thus threads must call back periodically

↳ yield() = thread voluntarily gives up on CPU - you ver → context switch

↳ preemptive threads = thread library tells OS to send it a signal periodically (timer)

↳ signal → signal handler → thread library → context switch

kernel level threads = OS can assign priorities to threads, but requires system calls, slower

- in user level threads if one thread blocks, entire process stops/blocked, can't use multiple CPU

posix threads (pthreads) interface =

```
/* hello.c - Pthreads "hello, world" program */
#include "csapp.h"

void *thread(void *vargp);

int main() {
    pthread_t tid;

    pthread_create(&tid, NULL, thread, NULL);
    pthread_join(tid, NULL);
    exit(0);
}

/* thread routine */
void *thread(void *vargp) {
    printf("Hello, world!\n");
    return NULL;
}
```

Thread attributes (usually NULL)

Thread arguments (void *p)

assigns return value (void **p)

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void *threadFn( void *ptr ) {
    char *message;
    message = (char *) ptr;
    printf("%s\n", message);
    /* do whatever you want */
}

main() {
    pthread_t thread1, thread2;
    char *msg1 = "Thread 1"; char *msg2 = "Thread 2";
    int iret1, iret2;

    /* Create independent threads each of which will execute function */
    iret1 = pthread_create( &thread1, NULL, threadFn, (void*) msg1);
    iret2 = pthread_create( &thread2, NULL, threadFn, (void*) msg2);

    /* Wait till threads are complete before main continues. Unless we
    /* wait we run the risk of executing an exit which will terminate
    /* the process and all threads before the threads have completed. */
    pthread_join( thread1, NULL); pthread_join( thread2, NULL);
    printf("Thread 1 returns: %d, Thread 2 returns: %d\n", iret1, iret2);
    exit(0);
}
```

joinable threads: can be reaped and killed by other threads • must be reaped to free memory resources

detached threads: cannot be reaped or killed by other threads • resources are automatically reaped on termination

signals to threads = when a process receives a signal, all threads share it

- every thread has its own signal handling → they can react differently to signal
- but a signal is received only once, random thread will take it