


synchronization

shared resources = global variables, heap → dynamically-allocated data, (stack is local, not shared)

access control mechanisms = low level: locks

↳ high level: mutexes, semaphores, monitors, condition variables




→ if one thread/process reads before other one writes → **race condition**

- the result is non-deterministic
- can happen only in writing mode

critical section = portion of the code that accesses the shared variables

locks



mutual exclusion = only one thread can execute code in their critical section one at a time

↳ all others are forced to wait, when one leaves other can enter

- shared (global) lock variable = 0 or 1

spinlocking = check lock variable if it is 1 → loop to wait while (lock-var)

atomic = means the code cannot be interrupted during execution. ↳ lock() and unlock() must be atomic

→ if there is a context switch in here, both try to get in → **race condition again**

atomic lock/unlock =

↳ hardware solution: → disable interrupts = whole process cannot receive necessary interrupts

- even noncritical section threads are locked
- only implemented at kernel level and inefficient on a multiprocessor

↳ **atomic update instruction** = cmpxchg src, dest - compare and exchange

- large performance penalty

↳ software solution:

Peterson's Algorithm


```
int flag[2] = {0,0};
int turn = 0;

// Thread 0
void plock0() {
    int other = 1;
    flag[0] = 1;
    turn = 0;
    while (flag[other] && turn == other) {
        ;
    }
}


void punlock0() {
    flag[0] = 0;
}

// Thread 1
void plock1() {
    int other = 0;
    flag[1] = 1;
    turn = 1;
    while (flag[other] && turn == other) {
        ;
    }
}

void punlock1() {
    flag[1] = 0;
}
```



Both Thread 0 and Thread 1 are ready to enter critical section. Race to update turn! Thread 1 writes FIRST, Thread 0 writes second. Thread 0 unlocks and starts spinlocking! Thread 1 enters critical section. Thread 1 unlocks. Thread 0 stops spinlocking and enters critical section. Thread 0 unlocks.



Both Thread 0 and Thread 1 are ready to enter critical section. Race to update turn! Thread 0 writes FIRST, Thread 1 writes last. Thread 1 loses and starts spinlocking! Thread 0 enters critical section. Thread 0 unlocks. Thread 1 stops spinlocking and enters critical section. Thread 1 unlocks.

- works completely fine, but slow and need to know how many threads

Spinlocks problem ⇒ threads waiting for lock waste CPU, there is still context switch to them, but they do not execute their code just looping.

mutexes

blocking locks • inserting spinlocks threads into a wait queue

- do not context switch them unless critical section is unlocked
- wake one thread from wait queue
- * frees up CPU for other threads

limitations ⇒ limiting writings but not limiting readings cannot be implemented

↳ exchanging data between two processes

Semaphores


to allow one thread to write or any number of threads to read at a time

semaphore = shared counter that initialized at the beginning init(sem, n)

↳ cannot be accessed directly

wait() - down() - p() = wait for semaphore value to become > 0, then decrement it

signal() - up() - v() = increment semaphore value by 1 → always executed no waits



```
semaphore my_semaphore;
init(my_semaphore, 1);

int withdraw(account, amount) {
    wait(my_semaphore);
    balance = get_balance(account);
    balance -= amount;
    put_balance(account, balance);
    signal(my_semaphore);
    return balance;
}
```

critical section

```
struct semaphore {
    int val;
    threadlist L; // List of threads waiting for semaphore
};

init(semaphore s, int initval):
    s.val = initval;

wait(semaphore s): //Wait until > 0 then decrement
    if (s.val <= 0) {
        add this thread to S.L;
        block(this thread);
    }
    s.val = s.val - 1;
    return;

signal(semaphore s): //Increment and wake up next thread
    s.val = s.val + 1;
    if (s.L is nonempty) {
        remove a thread T from S.L;
        wakeup(T);
    }
```

init(), wait() and signal() must be atomic actions!

the producer/consumer problem =

producer

consumer



int count = 0;



```
Producer Thread
Producer() {
    int item;
    while (TRUE) {
        item = produce();
        if (count == N)
            sleep(1);
        insertitem(item);
        count = count + 1;
        if (count == 1)
            wakeup(consumer);
    }
}
```

```
Consumer Thread
Consumer() {
    int item;
    while (TRUE) {
        if (count == 0)
            sleep(1);
        item = removeitem();
        count = count - 1;
        if (count == N-1)
            wakeup(producer);
        eat(item);
    }
}
```

if context switch → error

- pushes item, waits if full

- pulls item, waits if empty

using binary semaphore

if empty is 0, block itself

semaphore mutex; init(mutex, 1);
semaphore full; init(full, 0);
// N: size of the buffer
semaphore empty; init(empty, N);

Producer Thread

```
Producer() {
    int item;
    while (TRUE) {
        item = produce();
        wait(empty);
        wait(mutex);
        insertitem(item);
        signal(mutex);
        signal(full);
    }
}
```

Consumer Thread

```
Consumer() {
    int item;
    while (TRUE) {
        wait(full);
        wait(mutex);
        item = removeitem();
        signal(mutex);
        signal(empty);
        eat(item);
    }
}
```

- * mutex must be 1 and empty/full must be available to execute code

readers/writers

```
semaphore mutex; init(mutex, 1);
semaphore write; init(write, 1);
int readcount = 0;
```

```
Writer Thread
Writer() {
    wait(write);
    do_write();
    signal(write);
}
```

Reader Thread

```
Reader() {
    wait(mutex);
    readcount++;
    if (readcount == 1) {
        wait(write);
    }
    signal(mutex);
    do_read();
    wait(mutex);
    readcount--;
    if (readcount == 0) {
        signal(write);
    }
    signal(mutex);
}
```

only one writer will be available at a time

↳ if one at the end → only one reader at a time

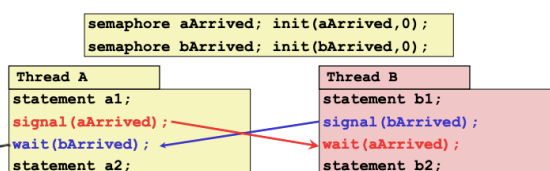
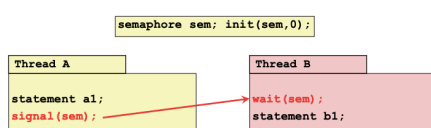
synchronization patterns

signalling = to guarantee that a1 is completed before b1 begins

rendezvous = same lines must be completed before next line

a1 must before b2, b1 must before a2

two A's cannot arrive because first one waits



barrier = generalized rendezvous, no thread executes critical point until after all the threads have executed rendezvous

```
n = numberofthreads;
count = 0;
semaphore mutex; init(mutex, 1);
semaphore barrier; init(barrier, 0);

All threads
rendezvous() {
    wait(mutex);
    count = count + 1;
    signal(mutex);
    if (count == n) signal(barrier);
    else {
        wait(barrier);
        signal(barrier);
    }
    criticalpoint();
}
```

→ can be used only once