

# file system implementation

MBR	partition table	partition 1	partition 2	partition 3	partition 4
-----	-----------------	-------------	-------------	-------------	-------------

\* each partition can have different file system & os

↓  
master boot record  
↓ first sector  
↓ loads boot loaders in partitions

outlines and describes the partitions

- EACH PARTITION -

boot block	super block	free space management	i-nodes	root directory	files and directories
------------	-------------	-----------------------	---------	----------------	-----------------------

loads os in kernel space

key info about file system which is read into memory at boot time

info about files (not in FAT)

## FAT (file allocation table)

- window's file system
- linked list implemented on an array
- stores free block list and file linkage

boot sector	reserved area	FAT1	FAT2	root directory	data area
-------------	---------------	------	------	----------------	-----------

↓ first sector

↓ vol name, # sectors, bytes, clusters...

optional copies of FAT1

↓ FAT12 and FAT16 only

FAT12 = 12 bits per entry →  $2^{12}$  entries →  $2^{12} \text{ bytes} = 8 \text{ KB}$  → max volume for 512B =  $2^9$  byte cluster  $2^{12} \cdot 2^9 = 2^{21} = 2 \text{ MB}$

FAT16 = 16 bits per entry →  $2^{16}$  entries →  $2^{16} \text{ bytes} = 128 \text{ KB}$

FAT32 = 28 bits per entry →  $2^{28}$  entries →  $2^{28} \text{ bytes} = 1 \text{ GB}$

FAT directory entries = name + attribute + disk block id

→ soft link for the pointers to same file

• FAT table entries are initially zero.

• -1 end of list (file A: 3 → 5 → 2 → -1)

## UFS (Unix File System)

boot area	super block	cylinder group 1	cylinder group 2	...	cylinder group N
-----------	-------------	------------------	------------------	-----	------------------

↓  
file system meta data

- EACH CYLINDER GROUP -

super block copy	cylinder group header	inode blocks	data blocks
------------------	-----------------------	--------------	-------------

- number of inodes and data blocks are fixed at file system creation

inode = each file has one, fits in a single block, all file attributes + pointers to data blocks

↳ data block pointers = direct, indirect, double-indirect, triple-indirect

for small data → they point to an array of data block pointers

↳ number of hard links = more than one directory can contain and point to same file, so its inode, only when the number is zero, delete the file

example ⇒

block size = 4K

block pointers = 4 byte

direct blocks = 12

direct = 12 block → 12

indirect = 4K/4 = 1K

double = 1K = 1024 block pointers

triple = ⇒ 1M blocks

triple = 1K \* 1K \* 1K ⇒ 1G blocks

max file

\* uses bitmap for empty blocks

file = inode + inode blocks + data blocks

directory entries = inode number + record length + name length + name + space

↳ basically inode number + name (no attributes)

↳ hard link for the pointers

## log structured file system - LFS

↳ other file systems update metadata frequently

↳ Lfs, do not overwrite blocks, instead each new write is appended at the end of the log segment

↳ imap keeps track of the last/valid version of the inodes

reason 1 = does not update the file in place every time (time consuming)

reason 2 = does not update the same block to prevent disk damage